

Memory Management

2026-03-10

1s2026

Gabriela Bittencourt, João Pedro Leôncio and Vinícius Peixoto

Contents



1	Understanding Resources	3
1.a	Memory x Storage	4
1.b	Program x Process	5
1.c	Memory Management Historical Advance	6
1.d	Non-contiguous Memory Allocation	7
2	Memory management overview	8
2.a	Physical vs. virtual address spaces	9
2.b	Some terminology	12
3	Physical memory allocation	13
3.a	Physical memory allocation	14
4	Virtual memory	16
4.a	Virtual memory	17
4.b	Page tables	18
4.c	Walking the page tables	21
5	References	22
5.a	References	23

1 Understanding Resources

Memory x Storage



When we talk about 'memory' in OS context, we are talking about **RAM**

Memory is NOT Storage

- Memory
 - fast, expensive, volatile
- Storage
 - slow, cheap, consistent

Program x Process



- Program
 - Code. Static entity.
 - Needs storage
 - Example: Your C code from MC202: the .c file
- Process
 - The execution of the code. Dynamic entity.
 - Needs memory: the instructions will be load in specific space of memory (), variables will be set to specific spaces in memory ()
 - Example: The instance of your text editor while you are writing a program; the instance of the terminal running when you are navigating through it to compile your code and when you run it; the instance of your program during the short period of time your computer is executing your code

Memory Management Historical Advance



- Contiguous Memory Allocation (no abstraction of memory)
 - ▶ One part is reserved to the OS, the other runs a single process
 - ▶ One process at the time: in order to change the process executed, the OS must save the previous process in disk and then load the next one.
 - ▶ Simple memory management
 - ▶ Poor memory utilization
- Non-contiguous Memory Allocation
 - ▶ Allows each process to “see” the entire address space as its own
 - ▶ Multiple processes can coexist
 - ▶ Memory allocation logic is opaque to processes – the kernel handles virtual memory

Non-contiguous Memory Allocation




Figure 1: Memory allocation of 5 processes (programs: A, B, C, D and A again) over time.

2 Memory management overview



Physical vs. virtual address spaces

- **Physical addresses:** actual addresses as seen by the CPU on the system memory bus (e.g. DDR RAM, memory-mapped PCIe registers, etc.)
- **Virtual addresses:** made-up memory addresses that are mapped onto physical memory addresses

Example: `vmlinux` (a Linux kernel image)

```
$ llvm-objdump -t ~/kernel/linux/vmlinux
/home/nuke/kernel/linux/vmlinux:      file format elf64-littleaarch64
SYMBOL TABLE:
ffff80008228b738 l      .init.text      0000000000000054 record_mmu_state
ffff80008228b790 l      .init.text      0000000000000030 preserve_boot_args
ffff8000822655c8 l      F .rodata.text 0000000000000034 __primary_switch
ffff80008228b738 l      .init.text      0000000000000000 $x
ffff80008228b7c0 l      F .init.text      000000000000008c __primary_switched
ffff800080061230 l      F .text      0000000000000020 set_cpu_boot_mode_flag
```


Note that the addresses for the symbols hover around address `0xffff8000822800` $\approx$ 65535 TiB (!)



Physical vs. virtual address spaces (ii)

- Clearly no system has that much memory
 - Might need to revisit this claim in the next decade or so :-)
- Those are **virtual addresses** configured by the kernel and translated into physical addresses
- This translation is done by a hardware unit called the Memory Management Unit (MMU)
 - The MMU is in turn configured by setting up data structures called **page tables**
 - Page tables describe the mapping between the virtual ↔ physical address spaces

Physical vs. virtual address spaces (iii)



CPU: Central Processing Unit

MMU: Memory Management Unit

TLB: Translation lookaside buffer

Figure 2: Overview of the address translation process.



Some terminology

- A **page** is a **block of memory** of fixed size
 - All memory is managed in units of pages
 - Traditionally 4KiB, more recently also 16KiB and 64KiB
 - In this class we always assume `PAGE_SIZE = 4096`
- Physical page number (**PPN**): the number of the n-th page in physical memory
 - Linux also calls this the **PFN** (page frame number)
 - Small exercise: what's the physical address of PPN=1024?
 - Another exercise: what's the PPN of physical address `0xffffffffabcde000`?
- Some commonly used idioms:
 - `PAGE_SHIFT`: number of bits to right-shift in order to get the PPN
 - For 4K pages, `PAGE_SHIFT = 12` ($2^{12} = 4096$)
 - `PAGE_SIZE`: the size of a page in bytes
 - `PAGE_SIZE = (1 << PAGE_SHIFT)`

3 Physical memory allocation

Physical memory allocation



Doing pretty much anything requires allocating physical memory. One of the first things a kernel needs to implement is a **physical page allocator**:

- Responsible for managing the available physical memory pages in the system
- Needs to keep track of which pages are in use and which ones are free
- Dumb (but perfectly valid) solution: freelist allocator
 - ▶ Linked list of all available pages
 - ▶ `alloc()` → pop the list head
 - ▶ `free(page)` → insert `page` at the front of the list




Figure 3: The freelist allocator (a linked-list of free pages)

Physical memory allocation (ii)



- This allocator has the upside of being dead simple to implement/debug
- But it would be nice to be able to allocate multiple contiguous pages at once
- In the next class we will go over other more elaborate allocators that solve this problem

4 Virtual memory

Virtual memory



- `riscv64` has 3 virtual memory modes (Sv39, Sv48 and Sv57)
 - We will work exclusively with Sv39 in this class
- We manage virtual memory by mapping **virtual page numbers** (VPNs) to **physical address numbers** (PPNs):




Figure 4: The Sv39 physical address encoding.




Figure 5: The Sv39 virtual address encoding.



Page tables

- In Sv39, a **page table** (PTB) is an array of $2^9 = 512$ words called **page table entries (PTEs)**
 - Note that in riscv64 1 word = 64 bits = 8 bytes and that $512 \times 8 \text{ bytes} = 4\text{KB}$
 - e.g. each page table takes up a physical page (not a coincidence)
- Each PTE can be either a:
 - **leaf page**: Map a VPN to a PPN
 - **non-leaf page**: Point to the PPN of another page table
- Sv39 uses a **3-level page table** scheme, which means that:
 - The **root page table** (also called a **level 1 page table**) contains entries that point to the PPN of a **level 2 page table**
 - Each **level 2 page table** points to a **level 3 page table** containing only leaf pages that map
- The process of following entries from root PTB $\rightarrow$ level 2 PTB $\rightarrow$ level 3 PTB until we reach a leaf page is called **walking** the page tables

Let's take a closer look at the structure of a PTE:

Page tables (ii)



63	62	61	60	54 53		28 27		19 18		10 9		8	7	6	5	4	3	2	1	0
N	PBMT	Reserved			PPN[2]		PPN[1]		PPN[0]		RSW		D	A	G	U	X	W	R	V
1	2	7			26		9		9		2		1	1	1	1	1	1	1	1

Figure 6: Sv39 page table entry.

The X/W/R bits have the following meaning:

X	W	R	meaning
0	0	0	Pointer to a next level page table
0	0	1	Read-only page
0	1	0	Reserved for future use
0	1	1	Read-write page
1	0	0	Execute-only page
1	0	1	Read-execute page
0	1	0	Reserved for future use
1	1	1	Read-write-execute page

Page tables (iii)





Figure 7: Sv39 page table entry.

For the others:

- **V** (valid): whether the PTE is valid
- **U** (user): whether the page is accessible in U-mode
- **G** (global): whether this is a global mapping that exists in all address spaces
- **A** (access): whether the page has been read/written/fetched since the bit was cleared
- **D** (dirty): whether the page has been written since the bit was cleared

Walking the page tables



We will go into more detail about how to actually implement this, but the (simplified) algorithm for walking the page tables is as follows:

- Start at the root page table (level 1)
- Check if the PTE at index `VPN[2]` has `V=1`; fail otherwise
- Move to the level 2 PTB at the physical address stored in `PPN[2:0]`
- Check if the PTE at index `VPN[2]` has `V=1`; fail otherwise
- Move to the level 1 PTB at the physical address stored in `PPN[2:0]`
- Extract the PPN that maps to the original VPN by looking at `PPN[2:0]`

5 References

References



- Tanenbaum, Andrew. (2012) '3. Memory Management', in *Modern operating systems*. Pearson Education, Inc., 4th ed.
- TODO: add kernel codes
- TODO: add kernel docs